

# NajNaf: WantCloud BV's Large Scale Image Resizer

F.W. Bakker\* S.F. van Wouw\*  
prof.dr.ir. D.H.J. Epema† dr.ir. A. Iosup† ir. B.I. Ghit‡

\*Authors {f.w.bakker,s.f.vanwouw}@student.tudelft.nl

†Instructors {d.h.j.epema,a.iosup}@tudelft.nl

‡Lab Assistant b.i.ghit@tudelft.nl

**Abstract**—WantCloud BV's current image resizing and persistent storage system is not able to cope with peak loads and does not scale up or down without human intervention. In this paper NajNaf is proposed: a system that is scalable, reliable, and durable without human intervention. The system is implemented as a cloud based application and has different subsystems that can be hosted in different virtual machines. This service oriented approach makes it possible to independently scale different parts of the system when necessary. Experiments with a prototype of the system on the DAS-4 cluster of the Delft University of Technology show that NajNaf dynamically scales up and down. This results in more than 90% of the tasks finishing before the global deadline. NajNaf ensures reliability by introducing redundancy into the system. It therefore is a viable alternative to the current system.



## 1 INTRODUCTION

WantCloud BV has an application that resizes and stores images in one relational database hosted by a shared webhosting provider. This system has one main issue: it does not scale automatically with the amount of concurrent users. Therefore the system cannot deal with peak loads efficiently. In addition, the relational database quickly becomes the bottleneck, because it has only one entry point.

In order to deal with the problems the current system has, a new system is proposed: NajNaf, of which the design is described in this paper. NajNaf is a large scale image resizer with persistent storage support that is specifically designed to be scalable, reliable, and durable without requiring human intervention. As such, this system deals with the problems WantCloud BV's current system has.

NajNaf can be hosted on a cloud infrastructure like Amazon EC2<sup>1</sup> or on a grid computing infrastructure like DAS-4<sup>2</sup>. The system has different subsystems which each can be hosted on a separate VM (Virtual Machine). The head node subsystem is responsible for monitoring all other subsystems, and scaling their capacity up or down when necessary. The master node is responsible for taking all image resize tasks and keeps track of them in a queue. The worker node subsystem consumes multiple image resize tasks off the master's queue and executes them parallel. Finally the persistent storage subsystem stores the resized images in a distributed object database.

If a subsystem fails it is automatically replaced by a

new VM running this subsystem. In case of the master this is a mirrored version (so the tasks in the queue do not get lost), and in case of a worker it is a new worker, after which the master automatically reschedules the lost image tasks.

Experiments with a prototype of the system (not implementing persistent storage) were conducted on the DAS-4 cluster using VMs with one cpu core available each. The experiments show that the system reacts to load changes and creates or suspends VMs accordingly within 15 seconds after the load change. This results in more than 90% of the tasks finishing before the global deadline.

The remainder of the paper is structured as follows: In section 2 requirements of the system are described, followed by the system design in section 3. The performance of the system is evaluated in section 4. Lastly, the pros and cons of the cloud based approach are evaluated in section 5 and conclusions are in 6.

## 2 REQUIREMENTS

There are two sets of requirements to the system: the requirements for the production system and the requirements for the prototype system. In this section we will first list the requirements of the prototype followed by the list of requirements for the production system. Note that the second list is an extension of the first list - all prototype requirements are also requirements for the production system. Functional requirements are annotated with *FR*, while non-functional requirements (constraints) are annotated with *NFR*.

**Prototype system:**

1. Amazon EC2 <http://aws.amazon.com/ec2/>

2. DAS-4 (The Distributed ASCI Supercomputer 4) <http://www.cs.vu.nl/das4/home.shtml>

**[FR:1] Image Resizing**

The system should resize the images such that they fit into the following dimension boxes (keeping aspect ratio intact):  $128 \times 128$ ,  $512 \times 512$  and  $1024 \times 1024$  (customizable by customer).

**[FR:2] Monitoring**

The system should include a monitoring-subsystem allowing administrators an insight into the performance of the system.

**[NFR:1] Automation**

Running the system should not require any human intervention.

**[NFR:2] Elasticity**

Upon a higher load the system's capacity must be automatically increased. When the load lowers the capacity must be reduced to minimize costs.

**[NFR:3] Load Balancing (Performance)**

The system should evenly distribute the tasks over the available VMs.

**[NFR:4] Modularity**

The system should have a Service Oriented Architecture (SOA) in order to be able to re-use different components within other WantCloud BV systems.

**[NFR:5] Reliability**

Downscaling the system (removing VMs) should not cause tasks to be dropped. Nor should tasks be dropped when a part of the system fails. Single point of failures are still allowed in the system (master and head node, see section 3).

**[NFR:6] Response time**

The time from when an image task is submitted until it has been completed should not exceed the global deadline of 30 seconds for at least 90% of all tasks.

**Production system:****[FR:3] Persistent Storage**

The resized images should be saved into a persistent storage facility.

**[NFR:7] Durability**

The persistent storage should make sure the stored images do not perish.

**[NFR:8] No single point of failure**

The system must be even more reliable by removing all single point of failures - all subsystems must support some form of replication.

**[NFR:9] Multi-tenancy Fairness**

The system must ensure the system is fair which is defined as: Users submitting first get served first and no starvation should occur (e.g. when one user submits a lot of jobs at once the users after this user should still get their jobs completed within finite time).

**3 SYSTEM DESIGN****3.1 Resource Management Architecture**

The global system design together with the allocation of system components to different types of VMs is depicted in figure 1. Components are grouped together on one VM type based on their coherence (as one would do when designing subsystems). However, components on the same VM type are decoupled as much as possible (using interfaces), such that they can be maintained independently ([NFR:4] Modularity). The system is described per VM type below. By referring to the requirements when explaining each component, one can infer whether the respective feature is actually implemented or not<sup>3</sup>.

**3.1.1 The Head Node**

The head node is responsible for managing all VMs (and itself is not a VM in our setup). It has a *VMMonitor* ([FR:2] Monitoring) which receives status updates from running VMs. The *VMScheduler* starts or stops VMs based on the feedback it receives from the *VMMonitor* ([NFR:2] Elasticity). This component is also responsible for the initial startup of the entire system ([NFR:1] Automation).

**3.1.2 The Master Node**

The system uses a master-worker setup to process images. There is one (possibly replicated, [NFR:8]) master, and there are multiple workers. A master node hosts the *TaskPool*, to which external producer programs (depicted by the External Producer VM type in figure 1) can push their image resize tasks. This *TaskPool* is implemented as a message queue with FIFO ordering (or can be modified to use a certain priority rule, depending on the fairness policy). The *TaskPoolMonitor* sends heartbeats to the *VMMonitor* containing the current resource utilization information and queue length. Finally the *MasterController* accepts commands from the *VMScheduler* (such as: prepare to shutdown).

**3.1.3 The Worker Nodes**

Multiple worker nodes are responsible for resizing images and uploading them to the persistent storage ([FR:1] Image Resizing). This is accomplished by a single *Consumer* per worker node that spawns a resize task thread for every task consumed from the *TaskPool*. When resizing has finished, the *StorageUploader* takes over and inserts the resized images into the persistent storage ([FR:3] Persistent Storage).

The *WorkerController* is similar to the *MasterController*. The *WorkerMonitor* reports the average image resize speed and the amount of concurrent image resize tasks it is running to the *VMMonitor*. The task of resizing

3. We have taken additional features such as durability and multi-tenancy fairness into account but did not actually implement them in the prototype version. We therefore did not include an *additional features* section. We consider the additional feature scheduling to be implemented by mechanic of section 3.2.5 and the master-worker setup

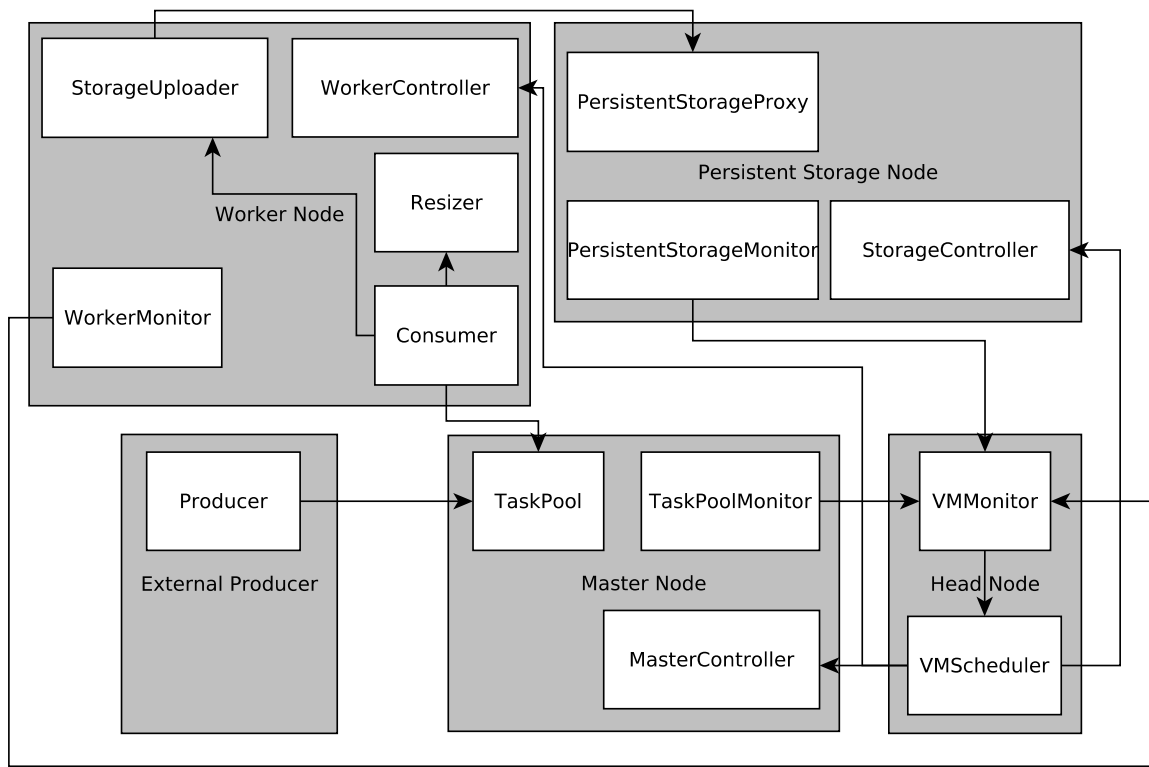


Fig. 1. System Design to VM Type Mapping

and uploading to storage is considered to be one atomic operation. The *Consumer* will only acknowledge back to the *TaskPool* if the entire task succeeded.

### 3.1.4 The Persistent Storage Nodes

Multiple persistent storage node VMs are responsible for storing the (resized) images ([FR:3] Persistent Storage). The *PersistentStorageProxy* is aware of the exact partitioning of persistent storage shards of a distributed object database such as CouchDB<sup>4</sup>. Replicating database clusters ensures [NFR:7] Durability.

The *PersistentStorageMonitor* and *StorageController* are based on the same idea as the *TaskPoolMonitor* and *MasterController*. The first sends heartbeats to the *VMMonitor*, and the latter accepts commands from the *VMScheduler*.

### 3.1.5 States and Commands

The worker and master's application state is used to inform the head node of the state the application is in (this state is important when the respective VM is actually up and running). Figure 2 shows the state diagram of the worker. The labels along the edges of the state diagram show what is required for the application to switch states. The 'err' is an internal event, while *INIT* and *SHUTDOWN* are commands received from the head node. A typical usage pattern is to send an *INIT* command with the master's ip address to a worker in

the *READY* state. And a *SHUTDOWN* command when the worker is no longer needed, after which the worker can be suspended if the worker enters the *READY* state again. If the head node detects a worker in the *ERROR* state, it will try to reset the worker by resending an *INIT* command. If this fails it simply deletes the worker.

## 3.2 System Policies

The system has different policies in place in order to comply with constraints set in section 2. These are explained in the sections below.

### 3.2.1 Scaling

The system is able to automatically scale up and down when needed ([NFR:2] Elasticity and [NFR:1] Automation). In the current design the persistent storage only scales up, and the worker nodes can scale up and down. The master node does not scale (it is assumed it can handle the queue on its own).

All VMs periodically send heartbeats with their current resource utilization to the *VMMonitor*. The monitor keeps track of the most recent heartbeats and periodically sends them to the *VMScheduler* (this means heartbeats can arrive in a certain timespan, resulting in a synchronised mechanism). The *VMScheduler* in turn will use the heartbeat data to determine what commands to issue (if any).

The *VMScheduler* will decide to scale up (add more worker nodes) if the total expected time it takes to

4. An example of how this can be accomplished with CouchDB can be found at <http://guide.couchdb.org/editions/1/en/clustering.html>

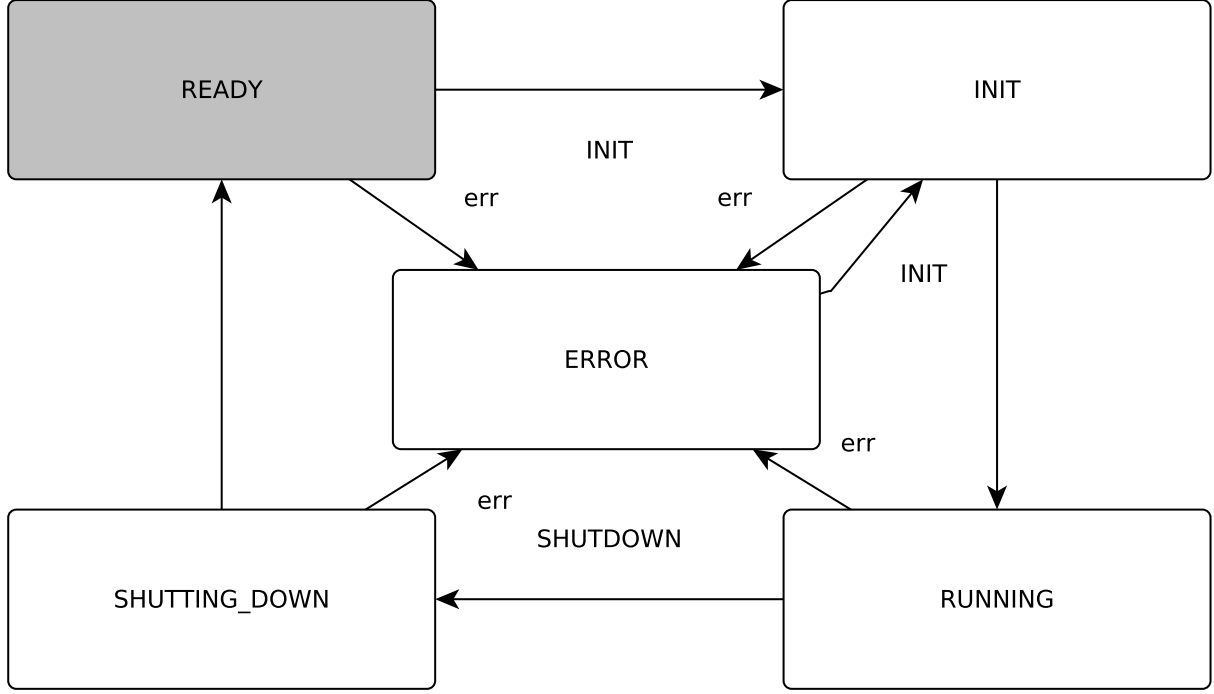


Fig. 2. State diagram of the worker

process all image tasks in the *TaskPool* is above a certain threshold (e.g. 30 seconds). It will scale down when the total expected time is below a certain threshold (e.g. 15 seconds)<sup>5</sup>. The amount of VMs that is created/resumed or suspended is determined by how far off track the total expected time is compared to the threshold. For example the current formula for scaling up is:

$$x = \left\lceil \frac{\text{expectedTime}}{\text{targetTime}} \right\rceil \cdot \text{workers} - \text{workers}$$

Assume the amount of VMs to add or remove to the pool of VMs is  $x$ , then in case of having to remove VMs from the pool, the *VMScheduler* will issue the *WorkControllers* of the  $x$  last added VMs to prepare for suspension. The corresponding *Consumers* will no longer consume new tasks from the *TaskPool* and the VM status will change from *RUNNING* to *READY*. The scheduler will then suspend the corresponding VMs when it will get to know the new status through a heartbeat.

Spawning new VMs can either be done by creating a new VM from a clean VM image or by resuming previously suspended VMs. The latter is preferred because initial research has shown that the time it takes to suspend and resume a VM is significantly faster than creating a new VM (about a factor 10). When a VM is

resumed and the size of the pool of suspended VMs is smaller than a certain threshold, new VMs are created and suspended. This helps maintain the size of the pool. When scaling down and the pool of suspended VMs will get above a certain threshold, VMs will be permanently removed from the pool instead<sup>6</sup>.

### 3.2.2 Load Balancing

The system needs to keep the tasks spread among all workers in order to comply with [NFR:3] Load Balancing (Performance), not overloading some nodes, while others are idle.

The first worker that is available to execute a task is the first worker that executes the task. The *Consumer* module on a worker node keeps consuming tasks from the *TaskPool* until it reaches the worker node's resource cap. This might lead to a load imbalance if some workers continuously get all the tasks while a couple of other workers suffer from starvation. But in this case the scaling policy kicks in, suspending the worker nodes that are under utilized. This results in fewer worker nodes being used, possibly increasing the load per node but also automatically making it more balanced.

Moving tasks between worker nodes would increase the complexity, but luckily there is no need for such a mechanism.

5. We actually started to let the *VMScheduler* base its decisions on the memory/disk and cpu loads of the workers. However, this did not take the amount of pending image tasks into account and cpu load averages were unreliable when VMs had just started.

6. This is currently not in the prototype, it will only create more VMs if necessary until it reached the VM cap. Then it only suspends and resumes VMs from its pool.

### 3.2.3 Reliability and Durability

In order for the system to be reliable and the persistent data to be really durable ([NFR:5] and [NFR:7] respectively), certain redundancy has been built into the system.

The persistent storage uses replication servers to replicate data to ensure durability. The master node VM running the *TaskPool* mirrors this pool on a different VM on a different physical machine (not implemented in the prototype). In addition, tasks are only acknowledged to the master node to be complete *iff* both the resizing and inserting into the persistent storage were successful.

If one of the VMs has not sent a heartbeat to the head node for at least 2 minutes (customizable) the *VMScheduler* will issue another VM of the same type to take over. In case of a master VM not responding, the scheduler will activate the mirrored VM to replace the master and instruct the worker VMs through their *WorkerController* to acknowledge their tasks to the new master VM instead. In case of a worker failing, the *VMScheduler* will simply spawn a new VM, and the *TaskPool* will reschedule the tasks that were lost due to not getting any acknowledgements. If one of the persistent storage nodes fails, the *VMScheduler* will promote a replica and creates one additional replica.

Finally, the head node could be replicated in future versions of the system to remove this single point of failure.

### 3.2.4 Fairness

The order of execution of tasks is determined by the order of the tasks in the *TaskPool*. In order to comply with [NFR:9] Multi-tenancy Fairness, we decided to use a FIFO queue implementation, which means tasks are executed in the order in which they are added. This has one potential downside however: if one user submits a large batch of tasks at once, future tasks may be delayed for a long time. Note that this will only be the case when the system is already scaled to its maximum capacity: if it hasn't yet the system will just scale up, preventing the users from suffering from the large batch of tasks. Although the fairness requirement is not violated in this scenario, one could implement more complex queues such as priority queues to solve this potential problem.

### 3.2.5 Slow Start Scheduling

The amount of images a *Worker* can resize at a given time depends on a few factors. The most important factor is the amount of processors the VM can use. Another factor would be the connection speed at which the *Worker* can download the image from the source. Whether the VM suffers from time slicing also effects how many threads can optimally run at the same time.

Due to the dynamic nature of two of these factors the worker has to dynamically determine how many image tasks (thus threads) it can process at the same time. It uses a special algorithm to achieve this. This algorithm

is based on the slow-start algorithm used by *TCP* [3]. This strategy is part of the congestion control algorithm of the *TCP* protocol. *TCP*'s slow-start uses a dynamic window size to determine how many packets can be sent at the same time - the *Worker*'s version uses this to determine how many threads it can run at the same time. *TCP*'s windows size is increased when an ack is received and reduced when an ack is not received in time, the *Worker*'s version decreases the window size when the VM's CPU is above a certain threshold, and increases the size when an image task has been processed without hitting this threshold.

## 4 EXPERIMENTAL RESULTS

### 4.1 Experimental Setup

NajNaf is written for the DAS-4 cluster using OpenNebula [2] for managing virtualization. Each VM runs a CentOS 5 image with 1GB of RAM and a single-core CPU of 2.4GHz (Virtual). All custom code is written in Python (version 2.7).

Communication between the VMs is accomplished with Python's Socket implementation. Memory and cpu usage of the VMs have been determined by using the unix `top` utility. The cpu usage is actually a load average that is determined by how many tasks want to use the cpu vs how many tasks have used the cpu in a certain timespan (1 minute in our case, to still be accurate but at the same time ignore huge cpu spikes). Disk usage was determined by the `df` unix utility.

The system consists of a couple of external libraries. These libraries are mentioned below.

**Apache-Apollo - an open source messaging queue system** Apache-Apollo is used to connect the producers (e.g. the image uploaders, or any other internal image creation process) with the consumers (the image resizers). It is essentially the *TaskPool* of the system. We decided to use a messaging queue because of the easy usage and because it creates a layer between the consumers and producers, allowing for easy and split scaling of the consumer and producer subsystems. Apache-Apollo was used as it supports the *STOMP* messaging protocol messaging protocol [4]. This has two benefits. First of all, when a new and better messaging queue system is created that supports *STOMP* it will be easier to replace Apache-Apollo. Secondly, the protocol is language independent, allowing the producers and consumers to be written in different programming languages. This makes the system more versatile.

**PIL - an open source image processing library for Python** One of the aspects of the system is the resizing and watermarking of the new images. For this, PIL is used. This library is able to perform all the required image manipulation tasks.

**LB\_watermarker** This simple library provides watermarking functionality [1] using PIL.

**CouchDB - image storage** The prototype does not yet contain an actual storage system. The eventual system

should contain a distributed and scalable data storage system, and we think CouchDB would be an interesting candidate due to its partitioning and scalability capabilities.

## 4.2 Experiment

A single experiment was executed. This experiment measures both the elasticity and the response time of the system. The parameters used in this experiment and the way the experiment was executed can be found in in section 4.2.1. The result can be found in section section 4.2.2.

### 4.2.1 Parameters and execution

The experiment required an artificial load. This workload was generated by adding  $i$  image tasks each  $j$  seconds to the Apollo messaging queue. The values of  $i$  were changed throughout the experiment. Even though the application consists of an actual download, resize and watermark script for resizing the image at a specific URL this was not used for the experiments. The main reasons are that it would have caused serious load on the server hosting the images and the performance would be dependent on the actual download times of the images, something that can vary throughout the experiment. Because of this the processing time of each image task was set to 2 – 3 seconds of busy waiting.

In order to not overload the DAS-4 cluster the size of the worker pool was set to 10 VMs. Note that the *VMScheduler* can easily manage many more nodes (the average load on the head node was neglectable). These workers were already created at the start of the algorithm.

Section 3.2.5 explains the slow-start algorithm the *Worker* uses. Due to the fact that each VM only has a single-core CPU available and the processing of an image task is just busy waiting, the worker always had just a single thread processing at a time.

The *VMScheduler* has a couple of configuration parameters that can be tweaked. One of these is the interval between the iterations in which decisions are made. During the experiments this was set to 15 seconds. As a result it takes at most thirty seconds before a worker VM starts working (15s to detect, 15s to detect running and issue an `INIT` command). The *VMScheduler* calculates the expected completion time - the time it will take all active workers to empty the queue. The scheduler attempts to keep this expected completion time between 15 and 30 seconds by either scaling up or down.

The experiment starts with a worker pool of 1 running and 9 suspended workers and a task pool of length 70. At the start of the experiment the queue is filled with 15 images tasks every 5 seconds. Each image task sent in this experiment causes the worker to busy wait for 2 – 3 seconds. At  $t = 200$  seconds this rate was decreased to just 10 image tasks every 5 seconds. At  $t = 870$  a peak load of 20 image tasks every 5 seconds was generated,

and at  $t = 930$  this load was changed to 5 image tasks every 5 seconds. At about  $t = 1200$  a peak load of 15 image tasks every 5 seconds was generated.

The charged-time and charged-cost of the experiments are evaluated by calculating the time the various VMs were running. We define the start of a run as the moment at which the VM was created or resumed, and the end of a run the moment at which a VM was suspended or deleted. The experiments are run on the DAS-4 cluster and not on Amazon EC2. Since the DAS-4 cluster does not charge for the usage of the cluster the charged-cost is taken as if the experiments were in fact run on Amazon EC2, assuming a cost of 10 Euro-cents/charged hour. Both the charged-time and charged-cost can be found at the end of section 4.2.2.

### 4.2.2 Results

According to ([NFR:2] Elasticity) it is essential that the system is able to scale both up and down automatically. The goal is to reduce the cost of running the system, but it is important that the response time of the system (the time between the image task entering the queue and the worker sending an *ack* to acknowledge it has finished processing the image task) remains within desired bounds. In this experiment the elasticity and response time is evaluated.

Figure 3 shows the expected completion time, the queue length and the number of worker threads during this experiment. Note that, due to the single-core the VM uses, each worker only had a single thread used to process image tasks. Thus, the amount of worker threads equals the amount of worker VMs. The figure shows a couple of interesting results. First of all, the amount of workers and the queue length follow the same pattern. This shows that, upon increasing the load, more workers are created in a timely manner. The expected completion time tends to fluctuate between 15 and 30 seconds, which is exactly the aim of the *VMScheduler*. There is only one occasion at which the expected completion time is significantly larger than 30 seconds, at  $t = 1300$ . This is due to the peak load of 20 image tasks generated every 5 seconds. The worker pool of 10 workers simply can't handle such a peak: 20 image tasks corresponds to a load of one minute of processing time. As these are generated every 5 seconds this will require a total of 12 workers to keep up. Note that even with this peak the average response time still satisfies [NFR:6] Response time.

Figure 4 shows the total expected sequential processing time (the time required to empty the queue if there had only been one worker) and the amount of worker threads. Both the amount of worker threads and the expected sequential processing time follow the same pattern, showing how well the system actually scales up and down.

The total sequential runtime of the experiment can be calculated by adding the runtimes of all workers and the master node. This gives a total of 14000 seconds, thus 3.9 charged-time hours. At a total of 10 Euro-cents/ charged

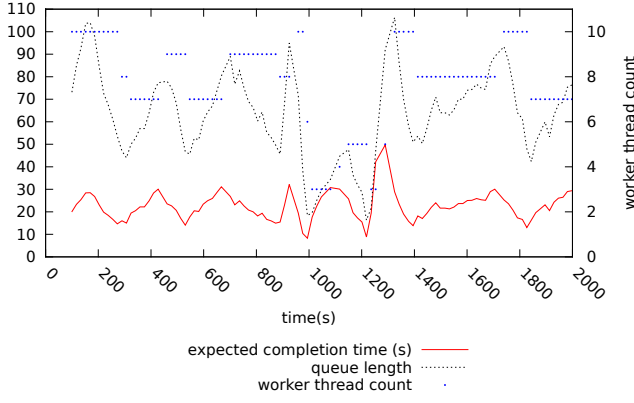


Fig. 3. Expected completion time, queue length and number of workers threads over time

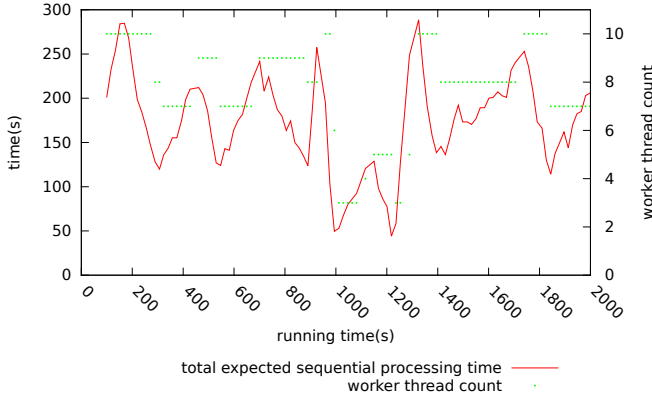


Fig. 4. Expected sequential processing time vs number of worker threads over time

hour the experiment cost a total of 39 cents. In this time about 4000 image tasks were processed, giving us a total cost of 0.01 cents per image task.

## 5 DISCUSSION

Although NajNaf is relatively complex for an image resize application because of the advanced requirements of the system, it proves to be a viable alternative to the current system WantCloud BV is using.

The experiments show that NajNaf satisfies both ([NFR:2] Elasticity) and ([NFR:6] Response time). The cost of processing one image is relatively low: section 4.2.2 shows the processing of a single image task (with a processing time of three seconds) costs as little as 0.01 Euro-cents. Since the system always needs exactly one active master node the cost of scaling a single image becomes slightly less as the size of the queue increases: the overhead of having the master node is spread out over more images. Thus when extrapolating over the found results, 10.000 images can be submitted for a total cost of 1 Euro, while the resizing of 10.000.000 images costs less than a 1000 Euro.

Apart from the requirements that are shown to be met by the experiments, there are several other features in

the prototype that make the prototype meet all of its requirements.

The head node keeps a log of all activities in the system. This log contains the latest states of the applications, the average processing times of the workers, the length of the messaging queue and many other statistics. The log itself is in human-readable format, satisfying ([FR:2] Monitoring).

Whenever a worker node crashes the head node deletes the VM and creates a new one. Since the worker only sends an acknowledgement after it has finished processing the image task it is guaranteed that no image tasks are lost. This satisfies both ([NFR:1] Automation) and ([NFR:5] Reliability). Whenever the master node crashes a new master node is created, and all workers are set to fetch image tasks from the new master. This also happens automatically, again conforming with ([NFR:1] Automation).

Due to the slow-start algorithm workers only fetch image tasks when they are able to actually process the image task. This conforms with ([NFR:3] Load Balancing (Performance)).

## 6 CONCLUSION

In this paper the design of NajNaf was described - a system that tries to deal with all the problems WantCloud BV's current system has. Experiments have been conducted on a prototype of the system and together with sound reasoning we have shown that that the system meets all the prototype requirements.

However, the prototype system does not contain all requirements, but provides a good base for the production system. All production requirements can be achieved by extending the prototype without exceptions. As a result we are convinced NajNaf can play an important part in the future of WantCloud BV. We therefore recommend WantCloud BV to choose for a cloud based solution, either off-the-shelf or custom made.

## REFERENCES

- [1] Lb\_watermarker. <http://www.turboradness.com/watermark-your-images-with-python/> [Visited 10 December 2012].
- [2] Opennebula. <http://opennebula.org/> [Visited 10 December 2012].
- [3] Slow start algorithm. <http://tools.ietf.org/html/rfc5681> [Visited 10 December 2012].
- [4] Stomp - the simple text oriented messaging protocol. <http://stomp.github.com/> [Visited 10 December 2012].

## APPENDIX A TIME SHEET

Activity	Time Spent (hour)
Thinking (design)	30
Developing	90
Experimenting	10
Analysis	10
Writing(report)	35
Wasted (setup etc.)	30
Total	205

The amount of time used for experimenting is fairly small compared to the other activities, this is mainly because the resource monitor was already in place and we only had to simulate workload for the system. Of the time spent on the experiments about 70% of the time was used for debugging and the remaining time was used to build a workload simulator and actually run the experiments.